

Machine Vision

Lecture Set – 04

Regions

Huei-Yung Lin

Robot Vision Lab

Regions and Edges

- Ideally, regions are bounded by closed contours
 - We could “fill” closed contours to obtain **regions**
 - We could “trace” regions to obtain **edges**
- Unfortunately, these procedures rarely produce satisfactory results



Region and Edges

- A **region** in an image is a group of connected pixels with similar properties
- Two approaches to partitioning an image into regions
 - *Region-based segmentation*
 - Value similarity and spatial proximity
 - Two pixels may be assigned to the same region if they have similar intensity values or if they are close to each other
 - *Boundary estimation using edge detection*
 - Edges are found based on differences between values of adjacent pixels
 - Most edge detectors use only intensity characteristics

Region and Edges

- In principle, region segmentation and edge detection should yield identical results
 - Edges may be obtained from regions using **boundary-following**
 - Regions may be obtained from edges using **region-filling**
 - (See the previous chapter)

Region Segmentation

- Region segmentation is a **partition** R_1, R_2, \dots, R_n such that
 - $\cup R_i = I$
 - $R_i \cap R_j = \emptyset$ if $i \neq j$
 - $P(R_i) = \text{TRUE}$
 - $P(R_i \cup R_j) = \text{FALSE}$, $i \neq j$, R_i adjacent R_j
- P is a function that evaluates similarities of the pixels in the region (or the **homogeneity predicate**)
- Gray-level to **binary image** conversion is a simple segmentation (partition to two sets)

3 / 30 / 2023

Segmentation Methods

- Thresholding
 - Global knowledge about an image or its part is usually represented by a histogram of image features
 - Also called histogram-based segmentation
- Edge-based
 - Based on the information about edges in the image
 - Use edge detecting operators
- Region-based
 - Construct the regions directly
 - Region-growing methods

Automatic Thresholding

- **Automatic thresholding** uses knowledge to select a proper threshold value
 - Intensity characteristics of objects
 - Sizes of the objects
 - Fractions of an image occupied by the objects
 - Number of different types of objects appearing in an image
- Automatic thresholding analyzes the gray value distribution in an image
 - Using a histogram of the gray values

Thresholding


- Gray-level thresholding is the simplest segmentation process
 - Objects or image regions are characterized by constant reflectivity of their surfaces
 - A brightness constant or threshold can be determined to segment objects and background
 - Computationally inexpensive and fast
- Thresholding is the transformation of an input image f to an output (segmented) binary image g as follows:

$$\begin{aligned}g(i, j) &= 1 \quad \text{for } f(i, j) \geq T \\ &= 0 \quad \text{for } f(i, j) \leq T\end{aligned}$$

where T is the threshold, $g(i, j) = 1$ for elements of objects, and $g(i, j) = 0$ for elements of background

Thresholding

- Algorithm: Basic thresholding
 - Search all the pixels $f(i, j)$ of the image f
 - An image element $g(i, j) = 1$ of the segmented image is an object pixel if $f(i, j) \geq T$, and is a background pixel otherwise

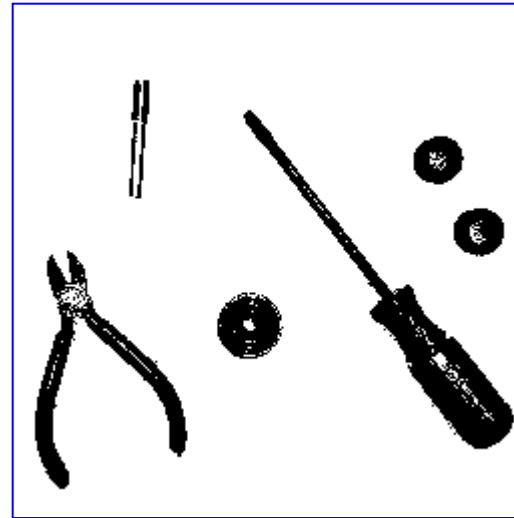


Example

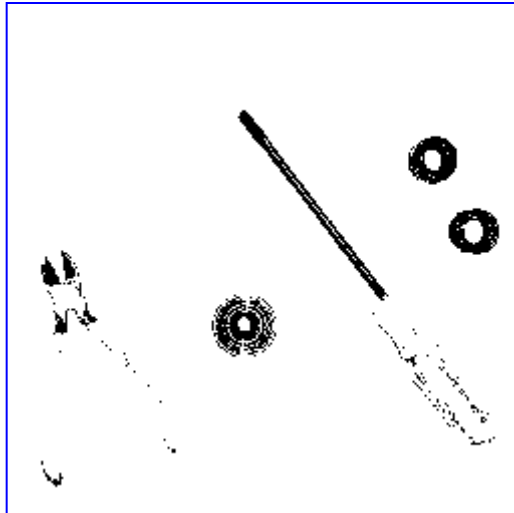
Original
image



Threshold
segmentation












Threshold
too low



Threshold
too high



Example




How to Find The Threshold Value?

- Manual
 - User defines a threshold
- P-Tile
- Mode
- Peakiness detection
- Iterative algorithm
- Other automatic methods


P-Tile Thresholding

- If we know the proportion of the image that is object (i.e., the **size** of the object is known!), threshold the image to select this proportion of pixels
- A printed text sheet may be an example if we know that characters of the text cover $1/p$ of the sheet area
- Choose a threshold T (based on the image histogram) such that $1/p$ of the image area has gray values less than T and the rest has gray values larger than T




Mode


- Based on histogram shape analysis
- Threshold at the minimum between the histogram's peaks




Example



Ideal histogram:




Add noise:



The valleys are good places for thresholding to separate regions

Finding the Peaks and Valleys


- Find a threshold T such that



- Methods:
 - Peakiness detection
 - Iterative threshold selection
 - Adaptive thresholding
 - Variable thresholding
 - Double thresholding

Peakiness Detection Algorithm

- Find the two highest **local** maxima in the histogram that are *at some minimum distance apart*
 - Suppose these occur at gray values g_i and g_j
- Find the lowest point g_k in the histogram H between g_i and g_j
- Find the **peakiness**, defined as $\min(H(g_i), H(g_j))/H(g_k)$
- Use the combination (g_i, g_j, g_k) with highest peakiness to threshold the image
- The value g_k is a good threshold to separate objects corresponding to g_i and g_j




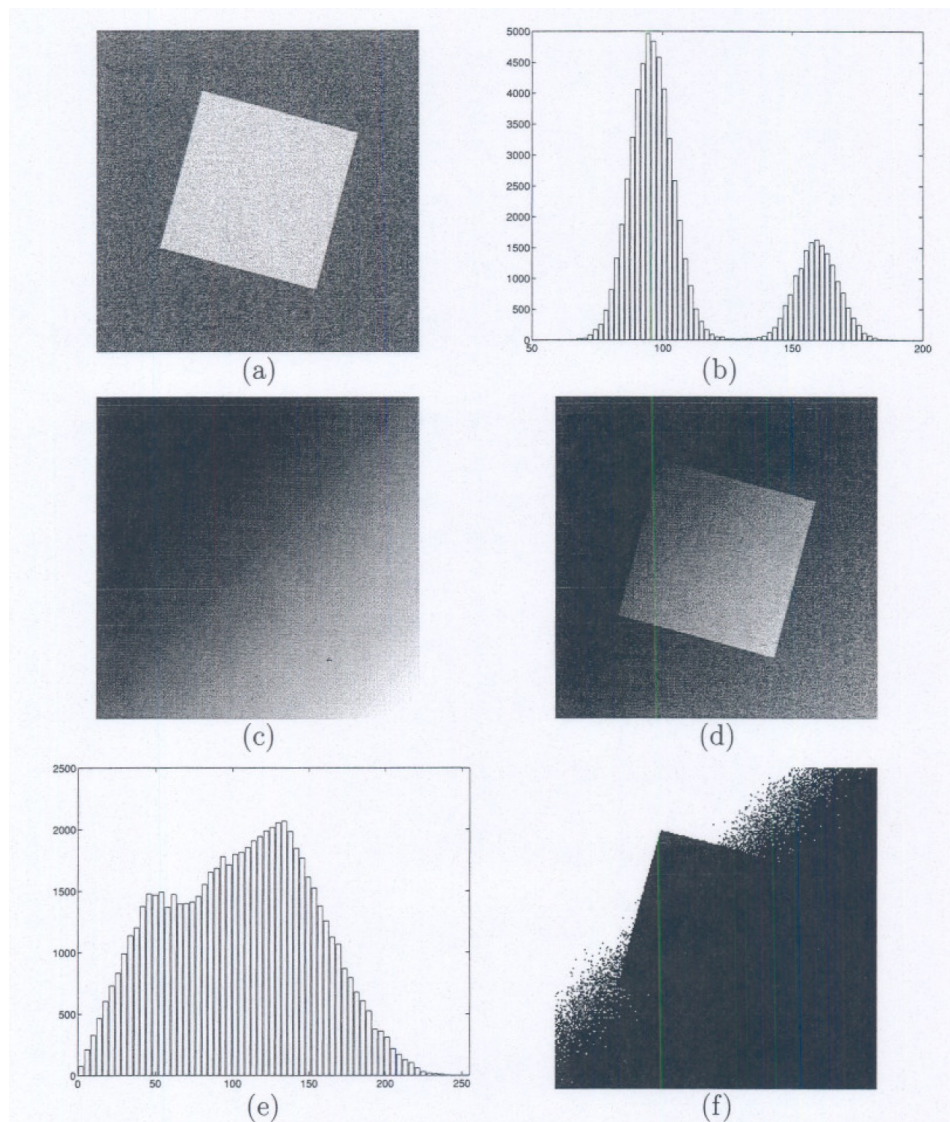
Iterative Threshold Selection

- Iterative threshold selection
 - It starts with an approximate threshold and then successively refines this estimate
- Algorithm: Iterative Threshold Selection
 - Select an initial estimate of the threshold T
A good initial value is the average intensity of the image
 - Partition image into two groups, R_1 and R_2 , using threshold T
 - Calculate mean gray values μ_1 and μ_2 of the partition R_1 and R_2
 - Select a new threshold $T = (\mu_1 + \mu_2)/2$
 - Repeat steps 2-4 until the mean value μ_1 and μ_2 converge

Adaptive Thresholding

- Automatic thresholding scheme may fail if the illumination in a scene is uneven
- The same threshold value may not be usable throughout the complete image
- Solution:
 - Partition the image into subimages
 - Select (different) threshold for each subimage based on its histogram
 - The final segmentation is the union of the regions of those subimages





Adaptive Thresholding



Variable Thresholding


- Also used to deal with **uneven illumination**
- Approximate the (new) intensity values of the image by a simple function such as a plane
 - The function fit is determined in large part by the gray value of the background
 - Histogram and threshold are then selected based on the fitted function
- Also called **background normalization**

Example



Double Thresholding


- To segment an object region which is inside another object region
 - Use a conservative threshold to obtain the core of the object and then grow the object regions
 - Use the second threshold to segment the object and background regions



Double Thresholding

- Double Thresholding for Region Growing
 - Select two thresholds T_1 and T_2
 - Partition the image into three regions:
 - R_1 : all pixels with gray values below T_1
 - R_2 : all pixels with gray values between T_1 and T_2
 - R_3 : all pixels with gray values above T_2
 - Visit each pixel assigned to region R_2
If the pixel has a neighbor in region R_1 , then reassign the pixel to region R_1
 - Repeat step 3 till no pixels are reassigned
 - Reassign any pixel left in region R_2 to region R_3

Example - Applications



Cross-Correlation Used
To Locate A Known
Target in an Image


Text Running
In Another
Direction

Limitations of Histogram Methods

- The most basic limitations
 - Use **global** information
 - Ignore **spatial** relationships among pixels
 - Use histogram, which contains no spatial information
 - Different patterns may have the same histogram!
- Other limitations
 - Constant illumination

Region Representation

- Classification of region representation
 - *Array representation*
 - *Hierarchical representation*
 - *Symbolic representation*
- Array representation
 - Use a same size array as the original image with entries indicating the region to which a pixel belongs
 - If an element of the array has value α , then the corresponding pixel in the image belongs to region α
 - For binary image
 - 0 : background
 - 1 : foreground



		1	1	1	1				
		1	1	1					
						2	2	2	
						2	2		
		3	3	3					
		3	3	3	3		4	4	4
		3	3	3	3		4	4	4
							4	4	4

Hierarchical Representation

- An image can be represented at different resolutions
 - *Lower resolution*: small array size, data loss
 - *Lower resolution*: less memory and computing requirements
- Hierarchical representation of images allows representation at multiple resolutions
 - Compute properties of image first at low resolution
 - Perform additional computations over a selected area at higher resolution
- Two different representations
 - “Pyramids” and “Quad trees”

Pyramids

- An **image pyramid** is a collection of representations of an image – the name comes from a visual analogy
 - Each layer of the pyramid is half the width and half the height of the previous layer (referred to as *coarse scale* version)
 - A pyramid is made by stacking the layers on top of each other
 - Pyramids are most convenient if the image dimension is power of two
- Common image pyramids
 - **Gaussian pyramid** (Gaussian * Gaussian = Gaussian)
 - **Laplacian pyramid**

4/6/2023

- HW #2 due TODAY
- HW #3 will be given next Thursday (one week late)
- Midterm on 4/20

The Gaussian Pyramid

- In a **Gaussian pyramid**, each layer is smoothed by a *symmetric Gaussian kernel* and resampled to get the next layer
- Let $S\downarrow$ denote downsample of an image, $P(I)_n$ denote the n th level of a pyramid $P(I)$, then


$$P_{\text{Gaussian}}(I)_{n+1} = S\downarrow (G_{\sigma} * P_{\text{Gaussian}}(I)_n) = S\downarrow G_{\sigma}(P_{\text{Gaussian}}(I)_n)$$

$$P_{\text{Gaussian}}(I)_1 = I$$

The Gaussian Pyramid

- The Gaussian pyramid is an example of image analysis by *a bank of filters* (smoothing filters)
 - It is a highly redundant representation because each layer is a low pass filtered version of the previous layer
 - we are representing the lowest spatial frequencies many times
 - A layer of the Gaussian is a prediction of the appearance of the next finer scale layer

The Gaussian Pyramid



The Laplacian Pyramid

- *A coarse layer of the Gaussian pyramid predicts the appearance of the next finer layer*
 - The coarsest scale layer of a **Laplacian pyramid** is the same as the coarsest scale of a Gaussian pyramid
 - Each of the finer scale layers of a L.P. is a difference between a layer of the G.P. and a prediction obtained by upsampling the next coarsest layer of the G.P.

$$P_{Laplacian}(I)_m = P_{Gaussian}(I)_m$$


$$P_{Laplacian}(I)_k = P_{Gaussian}(I)_k - S \uparrow (P_{Gaussian}(I)_{k+1})$$

The Laplacian Pyramid

- Each layer of the Laplacian pyramid can be thought of as the response of a *band-pass filter*
 - Low spatial frequencies being subtracted
- The L.P. can be used as an effective *image compression* scheme
 - Different level of the pyramid represent different spatial frequencies

The Laplacian Pyramid




- Each layer corresponds to the output of a band-pass filter
- The stripes give stronger responses at particular scales




Quadtrees

- Quadtrees
 - Trees where nodes have 4 children
 - Obtained by recursive splitting of an image
- Build quadtree:
 - Nodes represent regions
 - Every time a region is split, its node gives birth to 4 children
 - *Leaves* are nodes for uniform regions
- Merging:
 - Siblings that are “similar” can be merged

Quadtree Representation




not uniform



not uniform


1	2
3	4



uniform uniform

Quadtree for “Binary Image”

- Contains three types of nodes:
 - white, black and gray
- A region in an image is split into four subregions of identical size
 - If all points in that region are either black or white – no more splitting
 - If that region contains both black and white, called gray region – further splitting required
 - The splitting process continues until no gray regions



Symbolic Representation

- A region can be represented using symbolic characteristics
 - Enclosing rectangle
 - Centroid
 - Moments
 - Euler number
- Other non-symbolic representation
 - Mean and variance of intensities

Split and Merge

- After the initial intensity-based region segmentation, the regions need to be *refined* and *reformed*
- Why?
 - Simple intensity algorithms usually result in too many regions
- Due to
 - High frequency noise
 - Gradual transitions between regions
- Methods:
 - Interactively or automatically
 - May use domain and/or image processing knowledge

Split and Merge


- All pixels belong to region
 - Object
 - Part of object
 - Background
- Automatic refinement is done using a combination of split and merge operations
 - **Eliminate false boundaries** and spurious regions by merging adjacent regions that belong to the same object
 - **Add missing boundaries** by splitting regions that contain parts of different objects

Region Merging

- Merge **adjacent, similar** regions
- What does “similar” mean?
 - Examples:
 - “similar” average values : $|m_i - m_j| < T$
 - “small” spread of gray values: $|g_{\max} - g_{\min}| < T$
 - $g_{\max} = \max \{g(x,y) \mid (x,y) \in R_i \cup R_j\}$
 - $g_{\min} = \min \{g(x,y) \mid (x,y) \in R_i \cup R_j\}$
 - Other “similarity”:
 - luminance, statistical properties (mean, variance), color, texture, histogram shape, etc.
- Note:
 - A similar to B , and B similar to C does not imply that A is similar to C
 - Similar to some human behavior!!!


Region Adjacent Graph (RAG)

- A **region adjacent graph (RAG)** is used to represent regions and relationships among them in an image
- **Nodes** are used to represent regions, arc between nodes represent a common boundary between regions




Region Adjacent Graph (RAG)

- A **region adjacent graph (RAG)** is used to represent regions and relationships among them in an image.
- **Nodes** are used to represent regions, arc between nodes represent a common boundary between regions





Region Merging Algorithm

- From initial regions in the image using **thresholding** (on $n \times n$ regions, or manual selection) followed by **component labeling**





Region Merging Algorithm

- Prepare a RAG for the image
 - Regions are the **nodes**
 - Adjacency relations are the **links**





Region Merging Algorithm

- For each region in an image:
 - Consider its adjacent region and test to see if they are similar
 - Compare their mean intensities
 - Statistical characteristics (similar statistical distribution of intensity values)
 - Remove weak edges (common boundaries between regions)
 - If they are similar, merge them and modify the RAG



Region Merging Algorithm

- Repeat the previous step till no regions are merged



Using Similarity Measures

- Compare mean intensities (zero order model)
 - If the mean intensities do not differ by more than some predetermined value, the regions are considered similar
- Compare surface fitting (higher order model)
 - Model regions as polynomial surfaces
 - (e.g., take the gradual intensity changes into account)
 - Compute the fitting error
 - Merge regions if it decreases the fitting error

Using Statistics

- Model pixels as drawn from probability distributions
- Different regions have different distributions
- Assume the regions in an image have constant gray value corrupted by statistically independent, additive, zero-mean Gaussian noise
- Merge regions if their pixels came from the same distribution

Merge by Removing Weak Edges

- If the boundary between two regions is weak, then it can be removed and a large region can be obtained
- The “weak edge” can be specified by
 - *The intensity characteristics*
 - *The length of the common boundary*
- The common boundary is dissolved if the boundary is weak and the resulting boundary does not change gray value too quickly



Region Splitting

- If some property of a region is not constant, the region should be split
- Algorithm: **Region Splitting**
 - From initial regions in the image
 - For each region in an image, recursively perform the following steps:
 - Compute the *variance* in the gray value of the region
 - If the variance is above a threshold, split the region along the appropriate boundary
- Region splitting is generally more difficult than region merging

Split and Merge






- Split and merge are often used together
- All of the pixels in a region is homogeneous according to a predicate H , i.e., $H(R)$
- Suppose an image is partitioned into a set of regions $\{R_i\}$ and

$$\bigcup_{i=1}^n R_i = P \quad H(R_i) = \text{True} \quad H(R_i \cup R_j) = \text{False}$$



Split and Merge Algorithm

- Start with the entire image as a single region
- Pick a region R
 - If R is false, then split the region into four sub-regions
- Consider any two or more neighboring sub-regions, R_1, R_2, \dots, R_n , in the image
 - If $H(R_1 \cup R_2 \cup \dots \cup R_n)$ is true, merge the n regions into a single region
- Repeat these steps till no further splits or merges take place

Split



Merge



4/11/2023

- Midterm Exam next Thursday (4/20)
- No class next Tuesday (4/18)

Region Growing

- A set of algorithms to group pixels with similar attributes together
- The basic idea is to grow from a **seed pixel**
 - At a labeled pixel, check its neighbors
 - If the attributes of its neighbor is similar to the attributes of the labeled pixel, label the neighbor
 - Repeat until there is no pixel that can be labeled
- A simple case
 - The attribute of a pixel is its pixel value
 - The similarity given by the difference between two pixel values
 - If the difference is smaller than a threshold, then they are similar
 - Otherwise they are not


Region Growing


- Algorithm:
 - Generate initial seeds
 - Select a region
 - Add pixels to region
 - If growth stops, go back to step 2
- Efficient implementation
 - Based on **scan-line algorithm** in graphics
 - Each time we label a line instead of a pixel
 - This procedure is much more efficient than the recursive version


Watershed Segmentation

- Compare geographical watersheds
 - Divide landscape into catchment basins
- Edges correspond to watersheds
- Algorithm:
 - Locate local minima
 - Flood image from these points
 - When two floods meet
 - identify a watershed pixel
 - build a dam
 - continue flooding

Example












Watershed Point

Dam



Example



Marker-Controlled Watershed






Figure 5.51: *Watershed segmentation: (a) original; (b) gradient image, 3×3 Sobel edge detection, histogram equalized; (c) raw watershed segmentation; (d) watershed segmentation using region markers to control oversegmentation. Courtesy W. Higgins, Penn State University.*