

Machine Vision

Lecture Set – 06

Edge Detection

Huei-Yung Lin

What is an Image Feature?



What is an Image Feature?

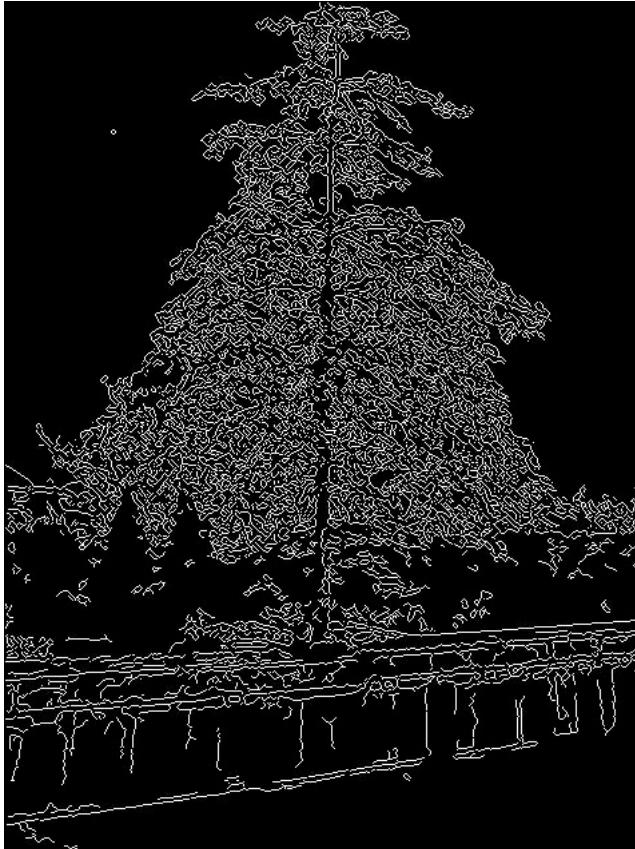


Image Features

■ Image features in computer vision

■ Global feature:

- A global property of an image
- e.g. the average grey value, the area in pixel, etc.

■ Local feature:

- A part of the image with some special properties
- e.g. a circle, a line, a textured regions, etc.



Image Features

- Image features (a practical definition): local meaningful detectable parts of the image
 - Points
 - Edges: step edges, line edges
 - Contours: closed contours are boundaries
 - Regions: similar color, similar feature, etc.

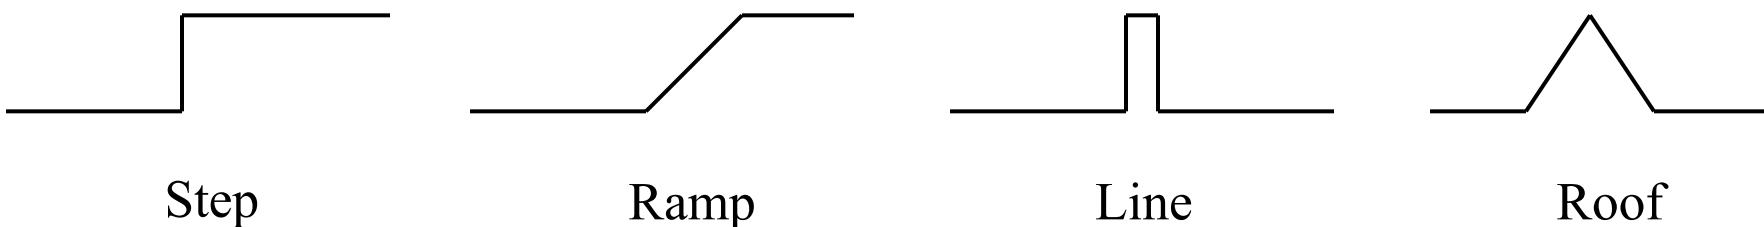


Remarks

- Most vision systems begin by detecting and locating some features in the input images
- In **3-D vision**, feature extraction is an *intermediate* step, not the goal
 - We do not extract lines just to obtain line maps
 - We extract lines to navigate robots in corridors or for camera calibration
- Feature extraction is for certain purpose of the system
- “Perfect” feature extraction is not necessary (and also not possible?)

Edge Detection

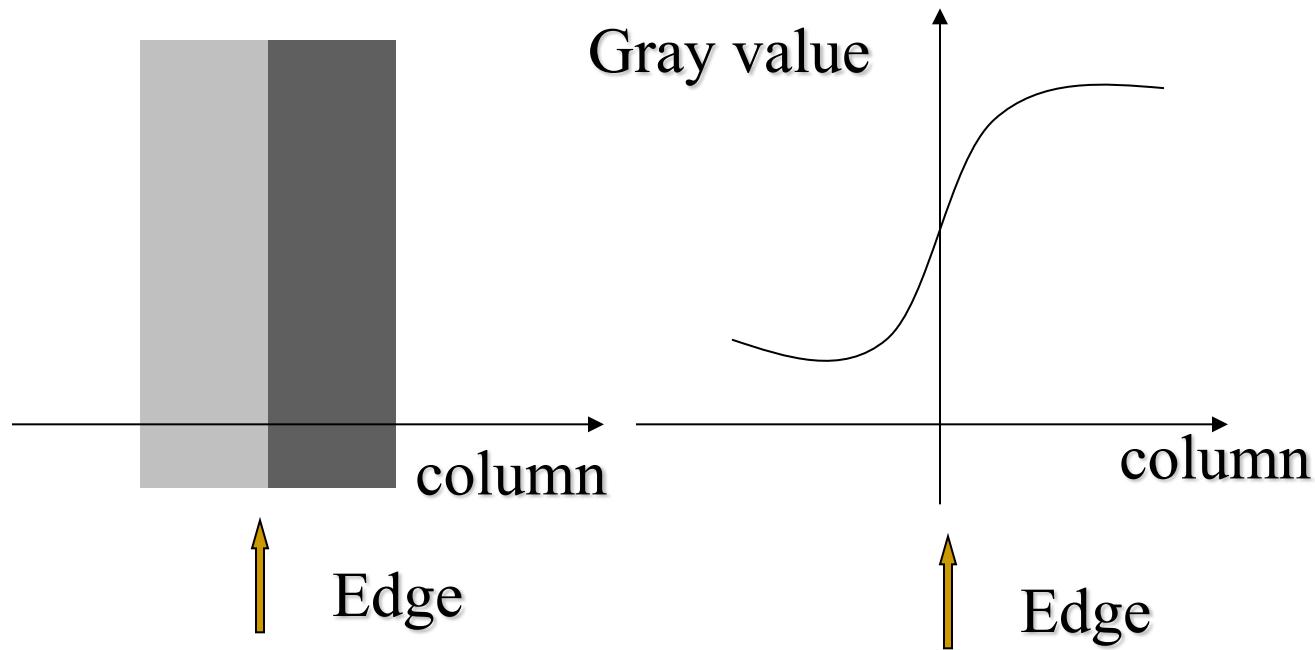
- Edges typically occur on the boundary between two different regions in an image
- Edge detection is frequently the first step in recovering information from images
- An edge is a significant local change in the intensity
 - Usually associated with a discontinuity in *either the image intensity or its first derivative*
 - The discontinuities can be step or line (ideally)
 - In reality, step becomes ramp and line becomes roof due to the smoothing of sharp edges in most images



5/2/2023

Edges

- They happen at places where the image values exhibit **sharp variation**



Definitions

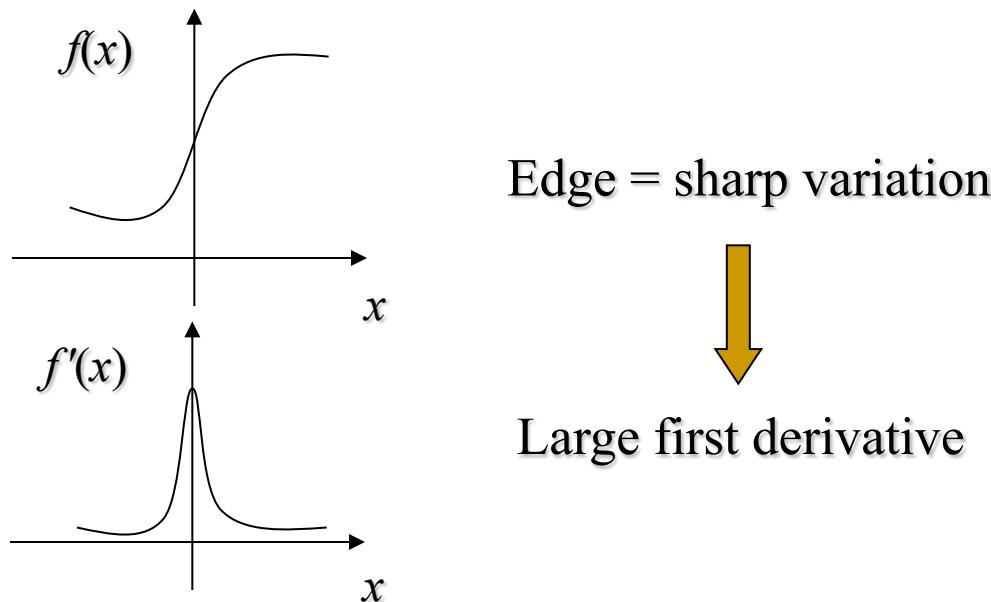
- An **edge point** is a point with coordinate $[i, j]$ at the location of a significant local intensity change
 - An **edge fragment** corresponds to the i and j coordinates of an edge and the **edge orientation** θ , which may be the **gradient angle**
 - An **edge detector** is an algorithm that produces a set of edges (edge points/edge fragments) from an image
 - A **contour** is a list of edges or the mathematical curve that models the list of edges
 - **Edge linking** is the process of forming an ordered list of edges from an *unordered list*
 - **Edge following** is the process of searching the image to determine contours
-

Edge Detection

- The term **edge** is commonly used for either **edge points** or **edge fragments**
 - An edge point may be the integer row and column indices of the pixel where the edge was detected or “at subpixel resolution”
 - An edge fragment may be conceptualized as a small line segment about the size of a pixel, or as a point with an orientation attribute
- There are three different kind of edges
 - Correct edges
 - False edges
 - Missing edges (true but not detected)

Edge Detection Scheme

- The basic approach to edge detection is to compute “spatial derivatives” of the intensity image
- The act of taking spatial derivatives is usually approximated by convolution



Gradient

- The **gradient** is a measure of change in a function
- Significant changes in gray values can be detected by using a “discrete approximation” to the gradient
- The gradient is a 2-D equivalent of the first derivate and is defined as a vector

$$\mathbf{G}[f(x, y)] = \begin{bmatrix} G_x \\ G_y \end{bmatrix} = \begin{bmatrix} \partial f / \partial x \\ \partial f / \partial y \end{bmatrix}$$

- The **vector** $\mathbf{G}[f(x, y)]$ points in the direction of the maximum rate of increase of the function $f(x, y)$
- The **magnitude** of the gradient given by $G[f(x, y)] = \sqrt{G_x^2 + G_y^2}$ equals the maximum rate of increase of $f(x, y)$ per unit distance

Approximation of Gradient

- The **absolute values** are commonly used to approximate the gradient magnitude

$$G[f(x, y)] \approx |G_x| + |G_y| \quad \text{or} \quad G[f(x, y)] \approx \max(|G_x|, |G_y|)$$

- The direction of gradient is define as

$$\alpha(x, y) = \tan^{-1}\left(\frac{G_y}{G_x}\right)$$

- The magnitude of the gradient is “independent” of the direction of the edge
- Such operators are called **isotropic operators**

Digital Approximation

- For digital images, the gradient approximation can be

$$G_x \approx f[i, j+1] - f[i, j] \quad \text{and} \quad G_y \approx f[i, j] - f[i+1, j]$$

- Simple convolution masks:

- 2×1 and 1×2

$$\begin{array}{|c|c|} \hline -1 & 1 \\ \hline \end{array} \quad \begin{array}{|c|} \hline 1 \\ \hline -1 \\ \hline \end{array}$$

- 2×2 and 2×2 (why?)

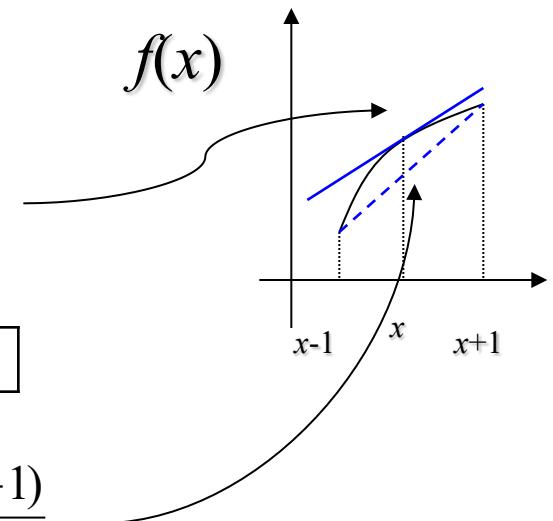
$$\begin{array}{|c|c|} \hline -1 & 1 \\ \hline -1 & 1 \\ \hline \end{array}$$

$$\frac{df(x)}{dx} = \lim_{\Delta x \rightarrow 0} \frac{f(x + \Delta x) - f(x)}{\Delta x}$$

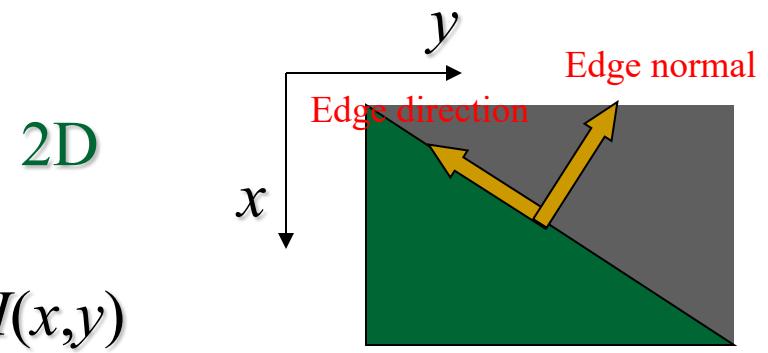
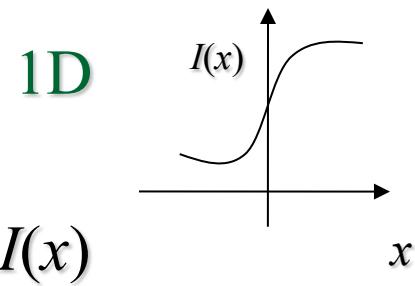
$$\begin{array}{|c|c|} \hline 1 & 1 \\ \hline -1 & -1 \\ \hline \end{array}$$

Convolve with: $\begin{array}{|c|c|c|} \hline -1 & 0 & 1 \\ \hline \end{array}$

$$\frac{df(x)}{dx} \approx \frac{f(x+1) - f(x-1)}{2}$$



Edge Detection



$$\frac{dI(x)}{dx}$$

$$\begin{aligned}\nabla I(x, y) &= [\partial I(x, y) / \partial x \quad \partial I(x, y) / \partial y]^T \\ &= [I_x(x, y) \quad I_y(x, y)]^T\end{aligned}$$

$$\left| \frac{dI(x)}{dx} \right| > threshold$$

$$\nabla I(x, y) = \sqrt{I_x(x, y)^2 + I_y(x, y)^2} > threshold$$

$$\tan \theta = \frac{I_x(x, y)}{I_y(x, y)}$$

Edge Detection

- Vertical edges:

- Convolve with:

-1	0	1
----	---	---

- Horizontal edges:

- Convolve with:

-1
0
1

The Essential Edge Descriptor

- Edge position or center
 - The image position at which the edge is located
 - Usually saved in a binary image (1 : edge, 0 : no edge)
- Edge normal
 - The direction (unit vector) of the maximum intensity variation at the edge point
 - This identifies the direction perpendicular to the edge
- Edge direction
 - The direction perpendicular to the edge normal
 - This identifies the direction tangent to the edge
- Edge strength
 - A measure of the local image contrast; i.e., how marked the intensity variation is across the edge along the normal

Edge Detection Steps

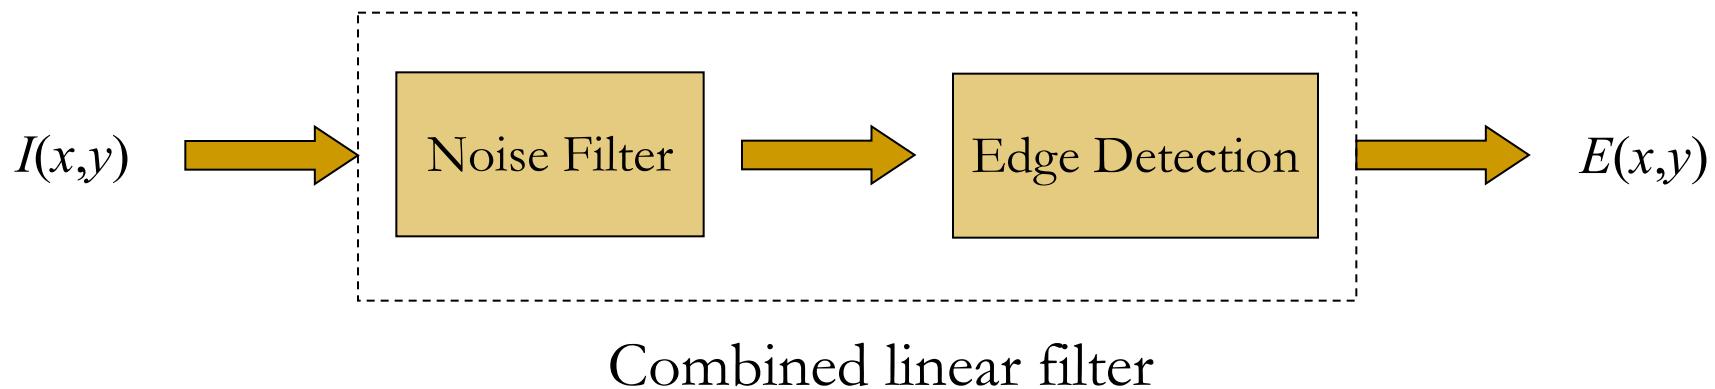
- Noise smoothing (filtering)
 - Suppress noise without destroying the true edges
- Edge enhancement
 - Design a filter responding to edges
 - Usually performed by computing gradient magnitude
- Edge localization
 - Thinning (non-maximum suppression)
 - Thresholding (used to decide whether the output is an edge point or not)

Some Assumptions

- The edge enhancement filter is linear
- The filter must be optimal for noisy step edge
- The image noise is **additive**, **white** and **Gaussian**
 - White noise: noise having a frequency spectrum that is continuous and uniform over a specified frequency band
 - White noise has equal power per hertz over the specified frequency band

Noise Suppression

- The **differential kernels** act as “high pass filters” which tend to amplify noise
- This is why edge detection is usually preceded by a noise reduction or filtering operation



Noise Smoothing & Edge Detection

■ Prewitt Edge Detector (vertical)

- Convolve with:

-1	0	1
-1	0	1
-1	0	1

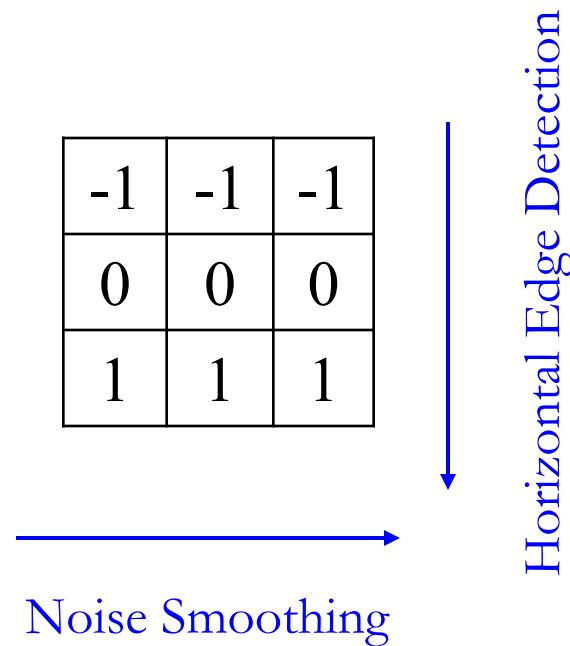
Noise Smoothing



Vertical Edge Detection

Noise Smoothing & Edge Detection

- Prewitt Edge Detector (horizontal)
 - Convolve with:



Roberts Detector

+1	0
0	-1

 G_x

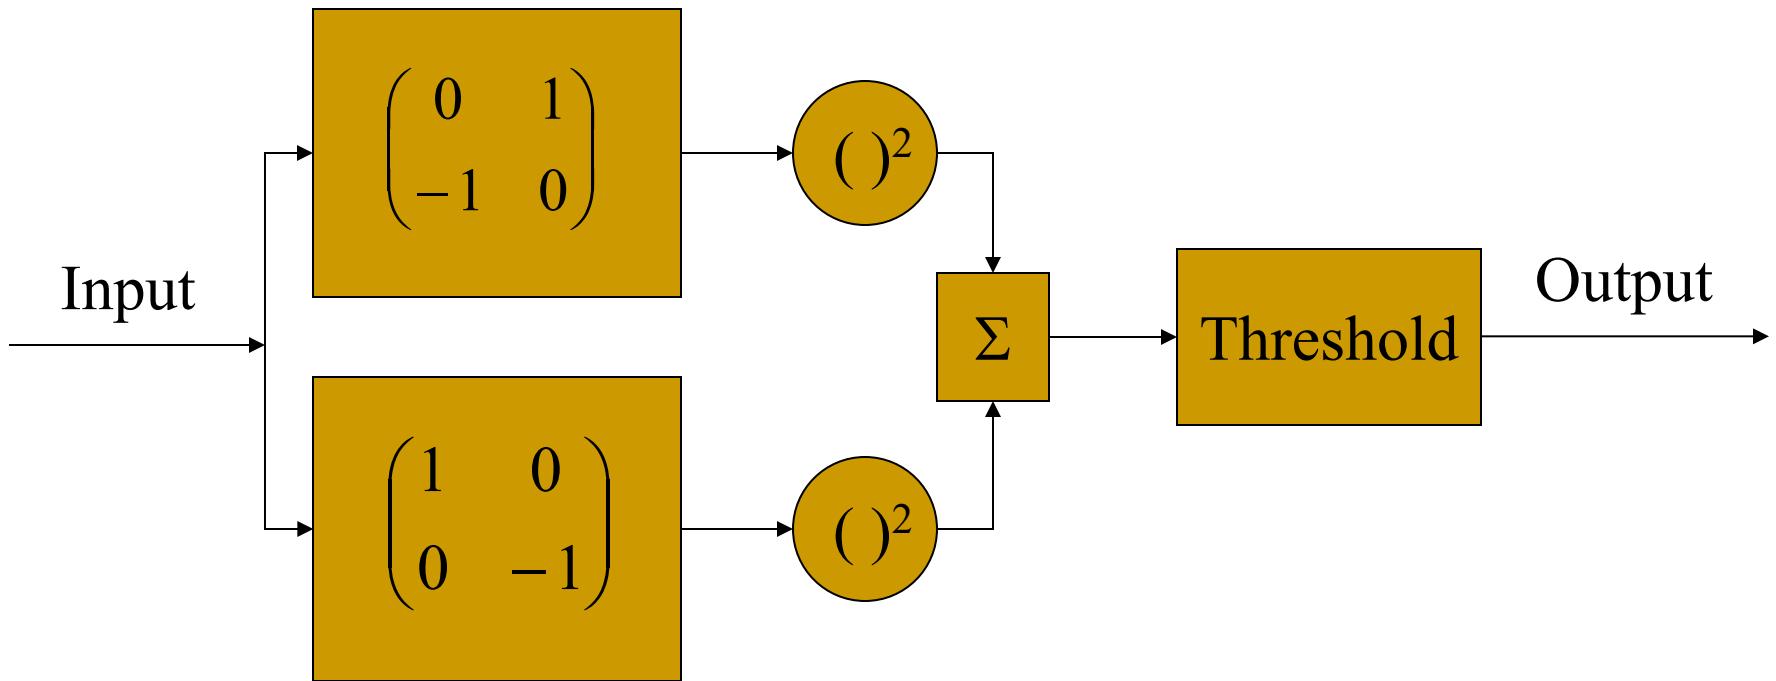
0	+1
-1	0

 G_y

$$|G| = \sqrt{{G_x}^2 + {G_y}^2}$$

Roberts Detector

- The Roberts detector gives an approximation to the continuous gradient at $[i+\frac{1}{2}, j+\frac{1}{2}]$. (not at $[i, j]$, why?)



Sobel Detector

- Sobel detector gives more weight to the 4-neighbors
- Emphasize the pixels closer to the center of the mask (compare with Prewitt!)

-1	0	1
-2	0	2
-1	0	1

 G_x

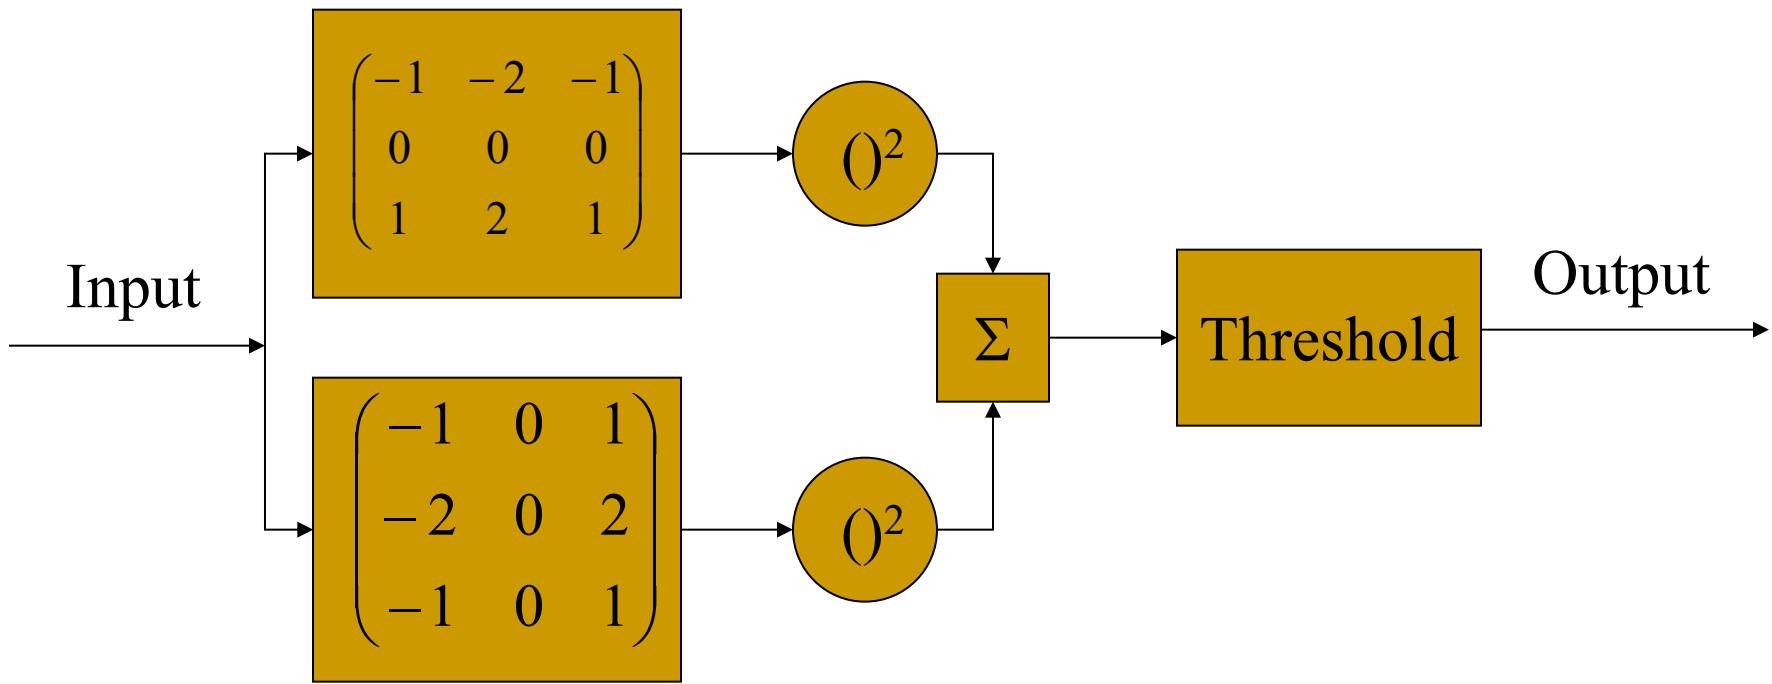
-1	-2	-1
0	0	0
1	2	1

 G_y

$$\theta = \tan^{-1}(G_y / G_x)$$

Sobel Detector

- Sobel operator is one of the most commonly used edge detector



Examples



Original



Sobel



Roberts

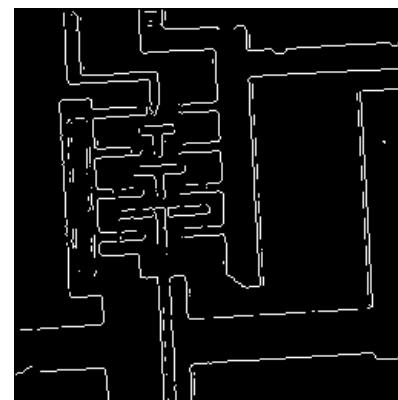
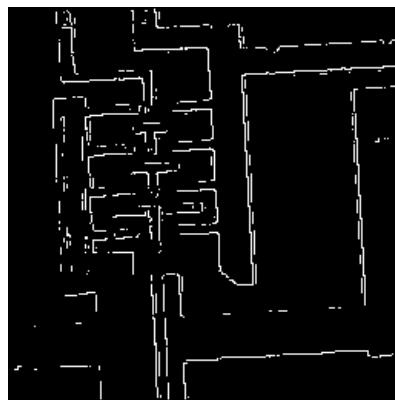
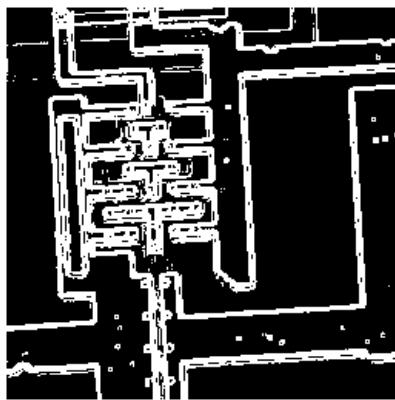
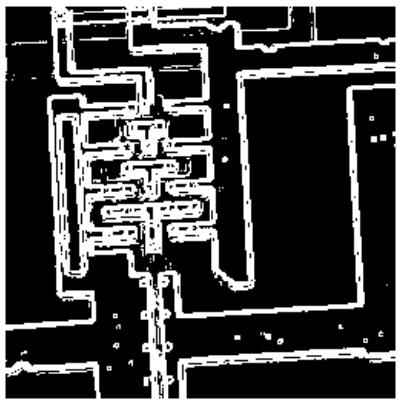
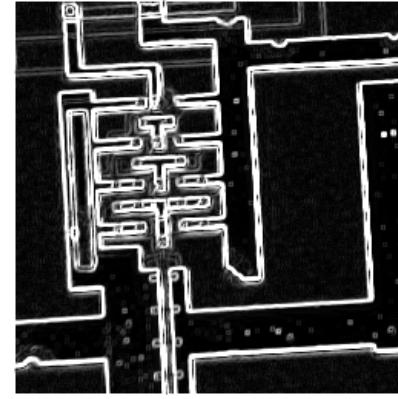
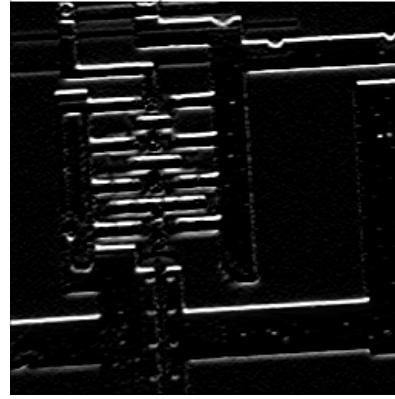
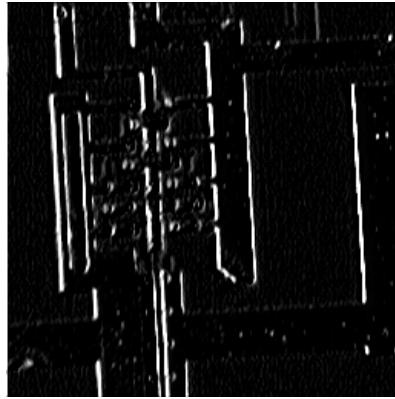
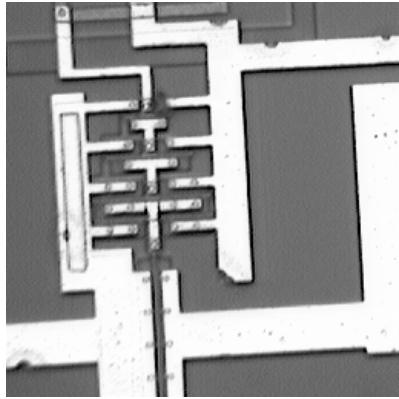


Canny



Prewitt

Example



Zero Crossing Detector

- Problem with **first-order derivative** edge detectors:
 - Using a threshold may result in detection of too many edge points
- Another approach:
 - Look for **zero crossings** of the **second derivative** since these will correspond to maxima (or minima) of the derivative
- $f''(x) = 0$ is not enough!
 - $f'(x) = c$ has $f''(x) = 0$, but there is no edge
- The second-derivative **must change sign**, i.e. from (+) to (-) or from (-) to (+)
- The sign transition depends on the intensity change of the image – i.e. from dark to bright or vice versa

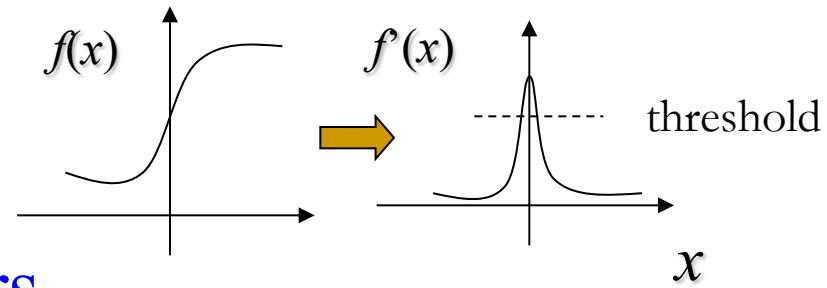
Zero Crossing Detector

- There are two operators in 2-D corresponding to second derivative
 - Laplacian
 - Second directional derivative

Edge Detectors

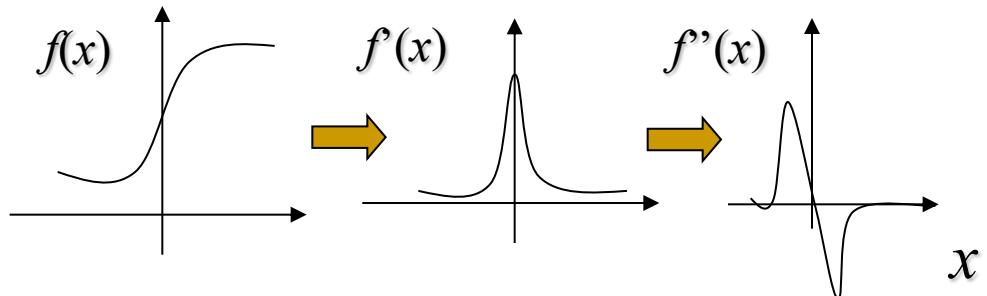
■ First-order derivative filters

- Sharp changes in gray level of the input image correspond to “peaks” of the first-derivative of the input signal



■ Second-order derivative filters

- Peaks of the first-derivative of the input signal, correspond to “zero-crossings” of the second-derivative of the input signal



Zero Crossings of the Laplacian

- The 2-D analog to the second derivative is called the Laplacian

$$\nabla^2 I = \frac{\partial^2 I}{\partial x^2} + \frac{\partial^2 I}{\partial y^2}$$



Laplacian

- Laplacian of f is defined as

$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

- Numerical approximation is given by

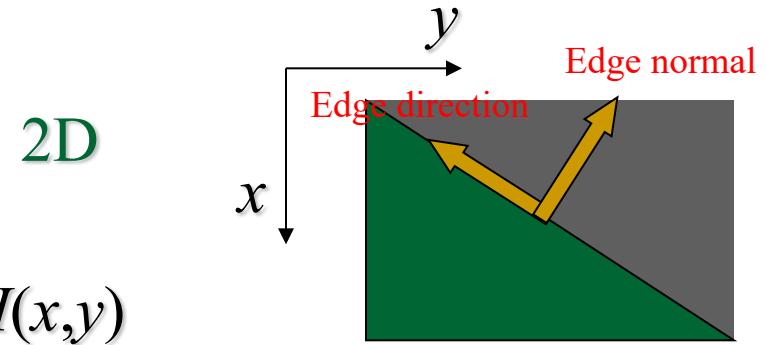
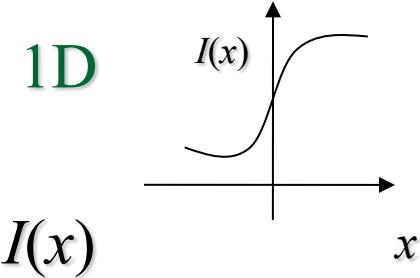
$$\frac{\partial^2 f}{\partial x^2} = f(i, j+1) - 2f(i, j) + f(i, j-1)$$

$$\frac{\partial^2 f}{\partial y^2} = f(i+1, j) - 2f(i, j) + f(i-1, j)$$

0	1	0
1	-4	1
0	1	0

1	4	1
4	-20	4
1	4	1

Laplacian Edge Detection



$$\frac{dI(x)}{dx}$$

$$\left| \frac{dI(x)}{dx} \right| > threshold$$

$$\frac{d^2I(x)}{dx^2} = 0$$

$$\begin{aligned}\nabla I(x, y) &= [\partial I(x, y) / \partial x \quad \partial I(x, y) / \partial y]^T \\ &= [I_x(x, y) \quad I_y(x, y)]^T\end{aligned}$$

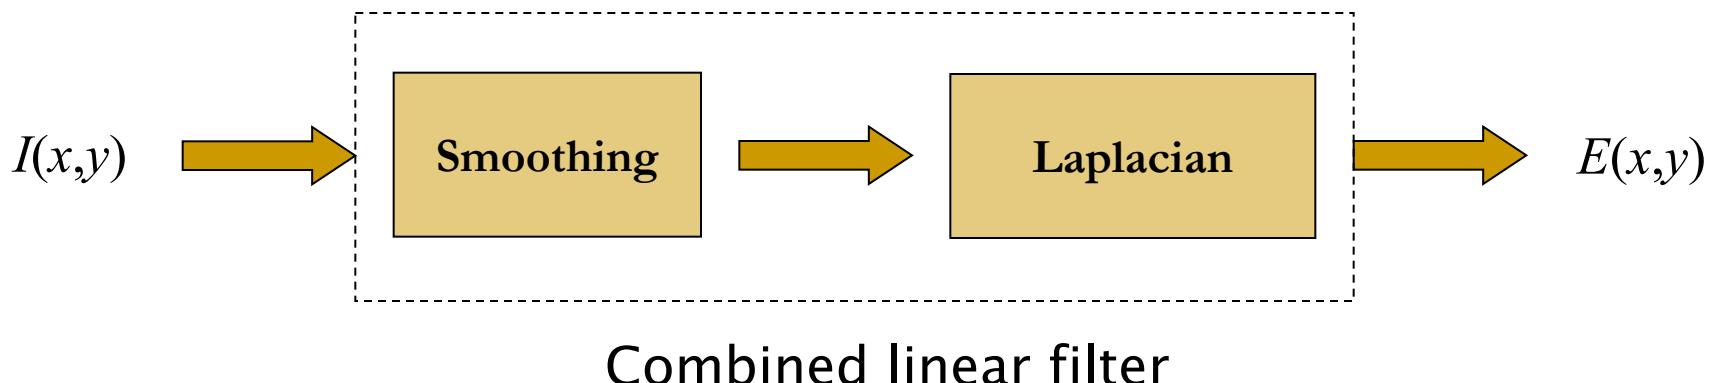
$$|\nabla I(x, y)| = \sqrt{I_x(x, y)^2 + I_y(x, y)^2} > threshold$$

$$\tan \theta = \frac{I_x(x, y)}{I_y(x, y)}$$

$$\nabla^2 I(x, y) = I_{xx}(x, y) + I_{yy}(x, y) = 0$$

Notes about Laplacian Filter

- $\nabla^2 I(x,y)$ is a **scalar**
 - Plus: Can be found using a “single” mask
 - Minus: Orientation information is lost
- $\nabla^2 I(x,y)$ is the sum of **second-order** derivatives
 - Taking derivatives increases noise
 - Very noise sensitive!
- It is “always” combined with a smoothing operation:



Approximating the Laplacian

- The second derivative along the x and y directions are approximated using difference equations*

$$\frac{\partial^2 f}{\partial x^2} = \frac{\partial G_x}{\partial x} = f[i, j+1] - 2f[i, j] + f[i, j-1]$$

$$\frac{\partial^2 f}{\partial y^2} = \frac{\partial G_y}{\partial y} = f[i+1, j] - 2f[i, j] + f[i-1, j]$$

- Combine the equations into a single operator, the Laplacian can be approximated by

$\nabla^2 :$

0	1	0
1	-4	1
0	1	0

Approximating the Laplacian

- Various discrete Laplacian kernels

0	1	0
1	-4	1
0	1	0

1	1	1
1	-8	1
1	1	1

-1	2	-1
2	-4	2
-1	2	-1

- For 1-D case:

1	-2	1
---	----	---

Zero-Crossing Edge Detector

- The zero-crossing corresponding to the edge could lie exactly halfway between the two pixels (the values, -7, 7, for instance)
 - The edge should be consistently marked at the pixel to the left or right of the edge
 - It rarely happens in most cases (usually -4, 6, etc.)
 - The actual edge location must be determined by interpolating the pixel values on either side of the zero-crossing
- The zero-crossing could correspond to a pixel in the image, but rarely happens (e.g. -3, 0, 3)

Laplacian of Gaussian (LoG)

- Before applying the Laplacian, the image is often pre-filtered with a Gaussian kernel
 - $O(x,y) = \nabla^2(G(x,y)*I(x,y))$
 - By linearity: $O(x,y) = \underline{\nabla^2 G(x,y)*I(x,y)}$
- Both stages can be combined into a single filter (The Mexican hat filter)

$$LoG(x,y) = -\frac{1}{\pi\sigma^4} \left[1 - \frac{x^2 + y^2}{2\sigma^2} \right] e^{-\frac{x^2+y^2}{2\sigma^2}}$$

1-D Gaussian

- 1-D Gaussian with zero mean is given by

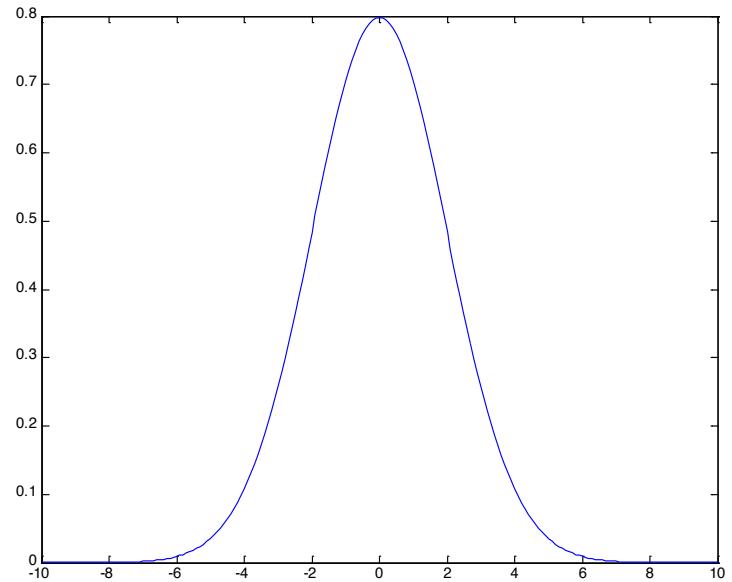
$$g(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{x^2}{2\sigma^2}}$$

- Ignore the scale factor,

$$g(x) = e^{-\frac{x^2}{2\sigma^2}}$$

- Then,

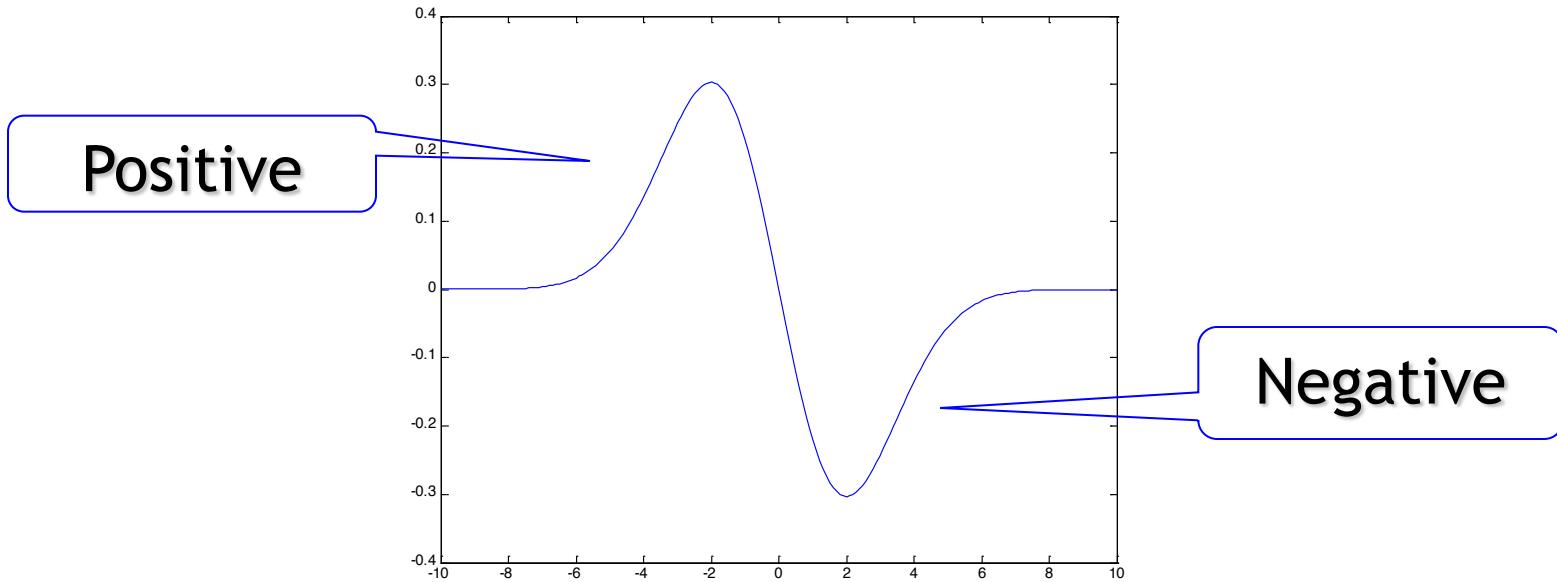
$$g'(x) = -\frac{1}{2\sigma^2} 2xe^{-\frac{x^2}{2\sigma^2}} = -\frac{x}{\sigma^2} e^{-\frac{x^2}{2\sigma^2}}$$



First Derivative of Gaussian

- The first derivative of Gaussian is given by

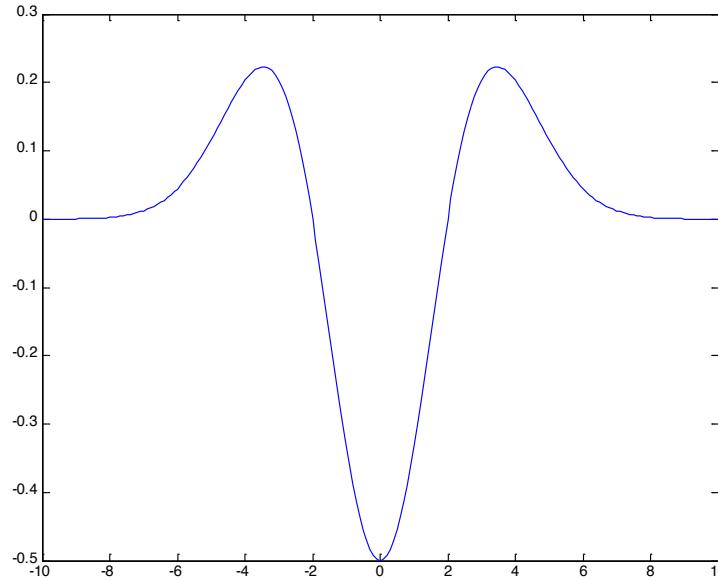
$$g'(x) = -\frac{1}{2\sigma^2} 2xe^{-\frac{x^2}{2\sigma^2}} = -\frac{x}{\sigma^2} e^{-\frac{x^2}{2\sigma^2}}$$



Second Derivative of Gaussian

- The second derivative of Gaussian is given by

$$g''(x) = \left(\frac{x^2}{\sigma^3} - \frac{1}{\sigma}\right)e^{-\frac{x^2}{2\sigma^2}}$$

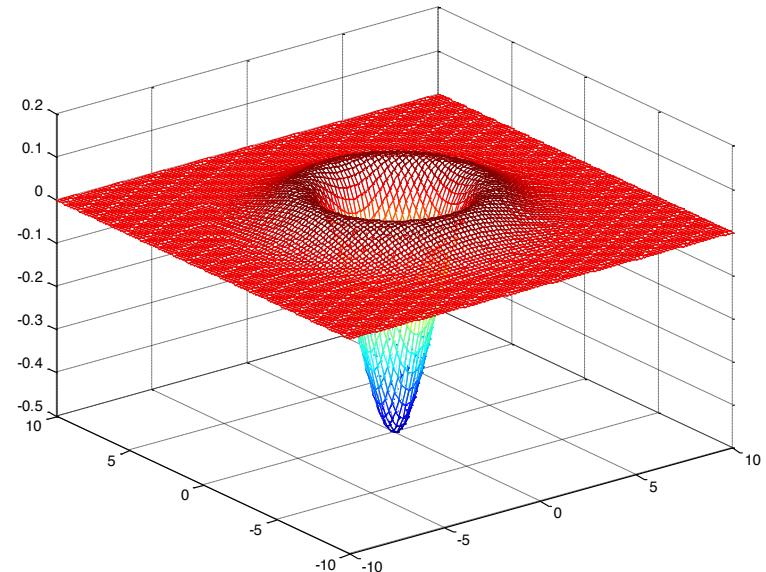


Laplacian of the Gaussian

- For the 2-D case, LoG is given by

$$\nabla^2 g(x, y) = \left(\frac{x^2 + y^2 - 2\sigma^2}{\sigma^4} \right) e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

“Mexican Hat”

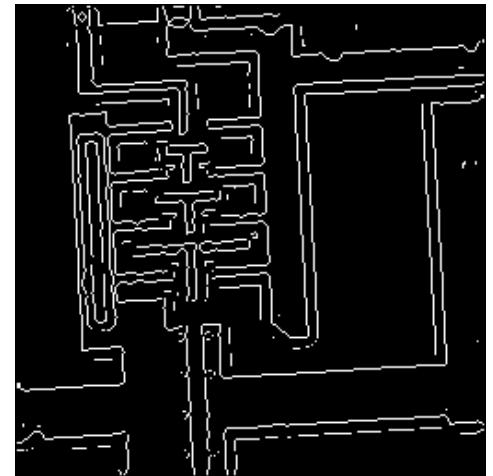
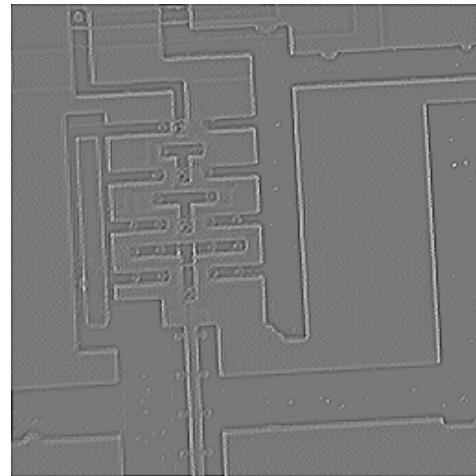
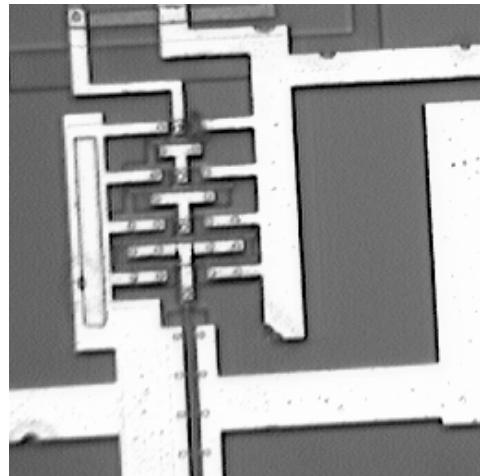


- LoG masks given in Fig. 5.12
- Scale space used in Gaussian smoothing
 - Ideal similar to double thresholding

LoG Edge Detector

- The fundamental characteristics of LoG
 - The smoothing filter is a Gaussian
 - The enhancement step is the second derivative (Laplacian in two dimensions)
 - The detection criterion is the presence of a zero crossing in the second derivative with a corresponding large peak in the first derivative
 - The edge location can be estimated with “subpixel” resolution using linear interpolation
- There are two ways of implementation
 - Apply Gaussian filter and compute the Laplacian
 - Apply Laplacian of Gaussian filter directly

Example



Disclaimer

- No edge detection scheme is going to work perfectly in all cases
- This is due to the fact that our notion of what constitutes a salient edge in the image is actually somewhat subtle

Detecting Corners

- We are often interested in detecting **point features** in an image
- These features are usually defined as regions in the image where there is significant edge strength in two or more directions
- They can be used for
 - Object tracking
 - 3D triangulation (stereo)
 - Object recognition

Detecting Corners

- Need two strong edges
- If E_x and E_y denote the gradients of the image in the x and y directions, then the behavior of the gradients in a region around a point can be obtained by considering the following matrix

$$C = \sum \begin{pmatrix} E_x \\ E_y \end{pmatrix} \begin{pmatrix} E_x & E_y \end{pmatrix} = \sum \begin{pmatrix} E_x^2 & E_x E_y \\ E_x E_y & E_y^2 \end{pmatrix}$$

Examining The Matrix

- One way to decide on the presence of a corner is to look at the eigenvalues of the 2 by 2 matrix C
 - If the area is a region of “constant intensity” we would expect both eigenvalues to be small (or zero)
 - If it contains a edge we expect one large eigenvalue and one small one
 - If it contains edges at two or more orientations we expect two large eigenvalues
- If $\min(\lambda_1, \lambda_2) > T$, then there is a corner!

Finding Corner

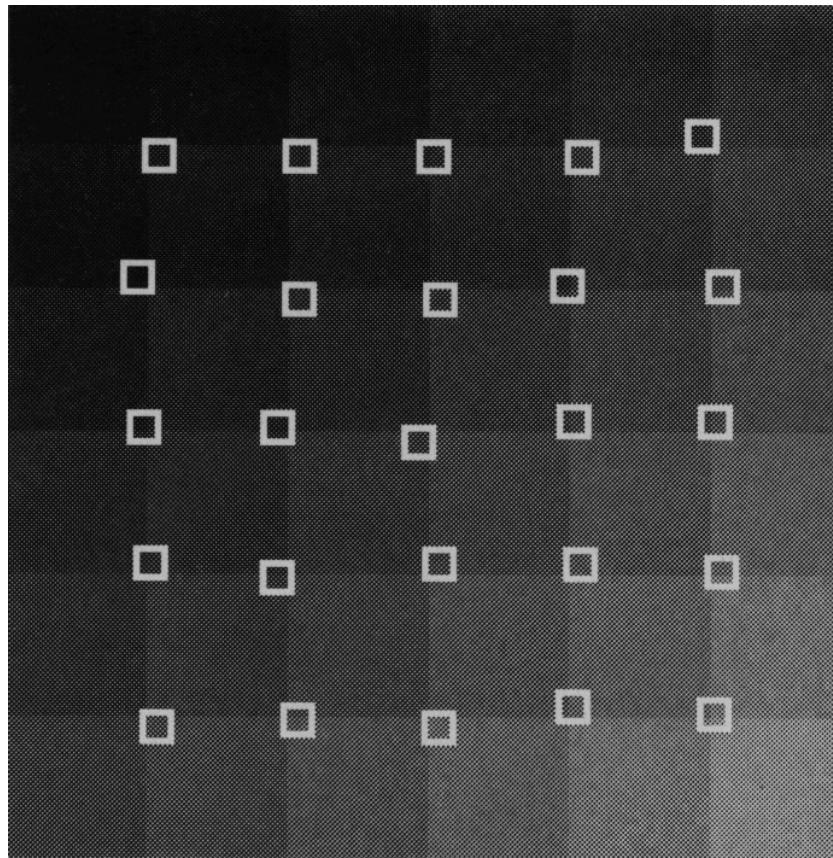
- One approach to finding corners is to find locations where the **smaller** eigenvalue is greater than some threshold
- We could also consider the **ratio** of the two eigenvalues
- Issues:
 - Localization – It can be difficult to precisely localize the corner in the intensity image
 - Modeling – It can be helpful to have a model of the corners you are trying to find in order to detect and localize them more systematically

Corner Detection

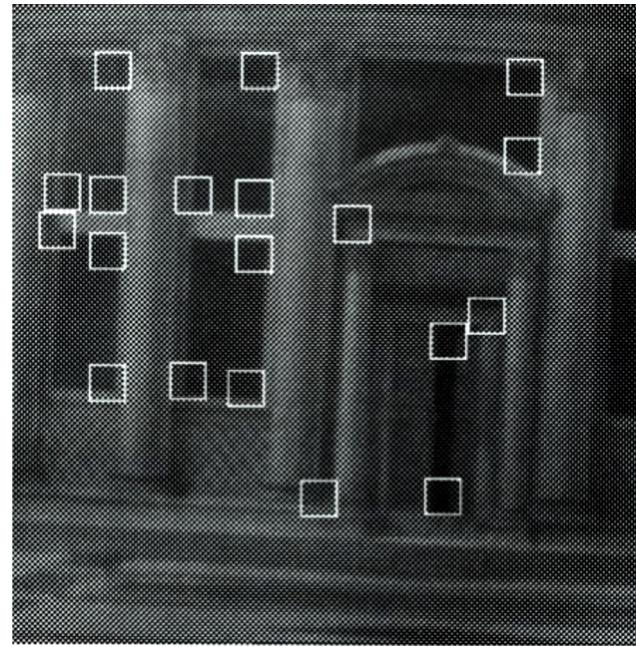
- Compute image gradient
- For each $m \times m$ neighborhood, compute matrix C
- If smaller eigenvalue λ_2 is greater than threshold τ , record a corner
- Nonmaximum suppression: only keep strongest corner in each $m \times m$ window

Corner Detection Results

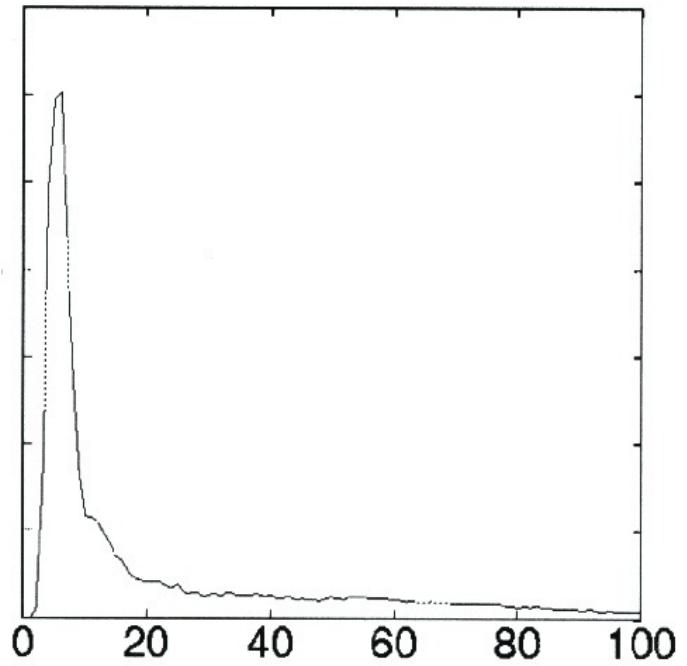
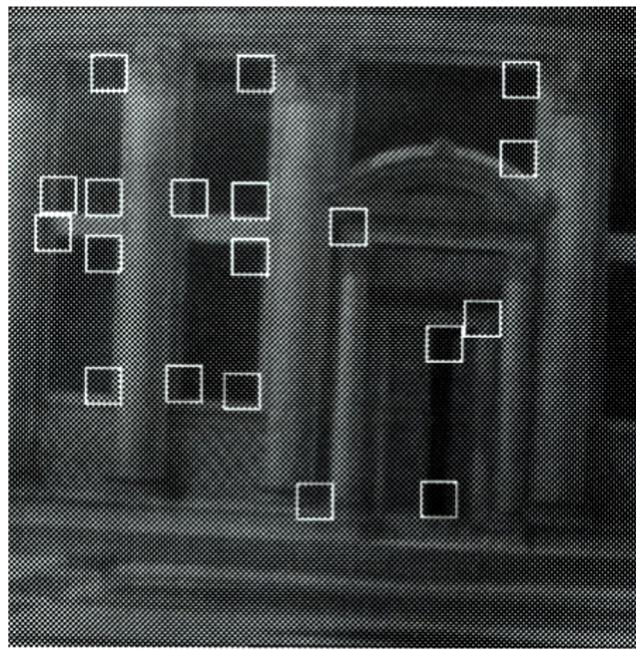
■ Checkerboard with noise



Corner Detection Results



Corner Detection Results



Histogram of λ_2 (smaller eigenvalue)